

# Prompt Iteration Log

Fabiana Alves Ferreira

*Corrected version. Fixes vs. the original: (1) Version 3 now tests exactly one layer — 'real context' — with a note explaining why the anti-invention sentence stays in the prompt without counting as a second layer; (2) every version has the four required notes; (3) Version 2 and Version 6 were relabeled to match exactly two of the ten allowed layers from the assignment — 'defined target audience' and 'verification requirements' — without changing what was actually run, just naming the tested layer correctly.*

## 1. Task chosen

Target task (FL-01): "Tailoring my resume for a specific job." Job used: Ab InBev — Data Quality Engineer (focus: Data Quality, ETL, SQL, Python, data testing). Base material: Fabiana's real resume and the Week 1 proof statement (Chile / TeleDesign).

## Standing project rule (not one of the tested layers)

Since Week 1, the instruction "never invent or exaggerate experience" has been a permanent rule for the whole project — not something created just for this exercise. That's why it appears in the prompt starting from Version 3 onward, but it doesn't count as the layer being tested in any specific version: it's a background constraint, present equally in every version after it.

---

## 2. Version 1 — Naive prompt (baseline)

**Prompt:** "Tweak my resume for this Data Quality Engineer job."

**Output (summary):** Generic advice disconnected from the actual resume: "add keywords like SQL, ETL and Python", "use action verbs", "be more concise". No sentence from the resume was rewritten; no reference to the real job posting or to Fabiana's story.

**What changed in the prompt:** Nothing yet — this is the baseline for comparison.

**What actually improved in the output:** Not applicable (baseline).

**What still didn't work:** The advice doesn't use the real resume, doesn't prioritize anything, and doesn't know what's already strong on the resume — it could have been written for anyone.

**What I'd try next:** Give the AI a specific role/audience so it evaluates like a specialist, not like a neutral assistant giving generic tips.

## 3. Version 2 — Defined target audience

**Prompt:** "You are a senior Data Quality Engineering recruiter with 10 years of experience screening junior resumes for ETL, SQL and Python roles. Evaluate the resume below for the Ab InBev Data Quality Engineer role..." [+ full resume pasted]

**Output (summary):** Specific, prioritized feedback: identifies that the line "ensured the quality and integrity of internal databases, eliminating duplicates" (Padma) is the strongest Data Quality signal on the resume, but it's buried mid-document. Recommends pulling it up. Flags that the Churn project (Python/SQL) should appear earlier. Flags that "data testing" doesn't explicitly appear in any listed experience.

**What changed in the prompt:** Defined the target audience the evaluation should be written for: a senior Data Quality Engineering recruiter with 10 years of experience screening junior resumes. In practice, this tells the AI whose reading perspective to evaluate from — the tested layer here is 'defined target audience'.

**What actually improved in the output:** Feedback stopped being generic and became prioritized — what to highlight vs. what's a weak signal — because the AI now evaluates the way someone who has screened hundreds of similar resumes would.

**What still didn't work:** None of the feedback was tied to my real story or my goal (getting an interview). There's a risk the AI tries to fill the 'data testing' gap with an invented term, like 'automated testing frameworks', which I've never done.

**What I'd try next:** Add real context — my proof statement and my goal — so suggestions tie back to something true, not just abstract resume optimization.

## 4. Version 3 — Real context

**Prompt:** Same prompt as V2, plus: "Context: I'm transitioning my career into data. My proof statement is: 'I identify inconsistencies in raw data and verify information against reliable sources before they cause rework — like when I validated illegible addresses for a network project in Chile, street by street, using Google Earth.' The goal is for a data manager to schedule a conversation with me. Never invent or exaggerate experience — every sentence must stay 100% true."

**Output (summary):** The AI started tying suggestions back to the proof statement: it recommended that the "eliminating duplicates" and "ensuring integrity" lines become the resume's throughline, not standalone items. It also explicitly refused to suggest "automated data testing frameworks" since there was no evidence of that in the resume, proposing instead the honest phrase "structured data validation via SQL and manual review".

**What changed in the prompt:** Added real context: my proof statement and my goal (having a data manager schedule a conversation with me). The tested layer here is real context — the anti-invention sentence was already present as a standing project rule, not a new variable in this version (see note above).

**What actually improved in the output:** Suggestions stopped being generic optimization and started pointing to a real throughline, consistent with who I am. The 'data testing' gap was named honestly instead of being filled with jargon.

**What still didn't work:** Context alone connects the suggestions to my story, but doesn't guarantee formatting consistency across the different experiences (Padma, Metaltex) — each one was still being rewritten in a different way.

**What I'd try next:** Give one full 'before and after' example for the AI to replicate the same pattern across the other experiences, instead of rewriting each one differently.

## 5. Version 4 — Examples (few-shot)

**Prompt:** Same context as V3, plus one full example: "Before: 'Ensured the quality and integrity of internal databases, eliminating duplicates.' After: 'Ensured the quality and integrity of internal databases, eliminating recurring duplicates and reducing operational rework — Data Quality discipline applied day to day.' Apply this same pattern... to the other experiences (Padma and Metaltex)."

**Output (summary):** All remaining Padma and Metaltex lines were rewritten following the same three-part mold. The result was consistent in format and tone — but repetitive: nearly every sentence ended with some variation of "Data Quality discipline applied to X", which reads mechanically if used without review.

**What changed in the prompt:** Added one full example (before/after) for the AI to replicate the pattern across the other experiences.

**What actually improved in the output:** The example acted as a mold and guaranteed structural consistency between sentences — it solved V3's tone inconsistency.

**What still didn't work:** The AI over-generalized the example's pattern, applying the same formula repeatedly. It was consistent, but mechanical — as it stood, I wouldn't use it without rewriting it again.

**What I'd try next:** Ask for a specific output structure (table) to make row-by-row review easier, and/or explicitly ask for sentence variation so it doesn't sound robotic.

## 6. Version 5 — Output format

**Prompt:** Same context as V4, plus: "Deliver the answer as a table with three columns: Original | Suggestion | Justification (max 20 words each). At the end, also rewrite the professional summary in 3 lines."

**Output (summary):** A clean, row-by-row table, each original sentence next to its suggestion and a short justification — easy to approve or reject item by item. The final professional summary came out as 3 focused lines, citing SQL, Python and Data Quality right at the start.

**What changed in the prompt:** Asked for a specific output format: a three-column table, plus a 3-line professional summary.

**What actually improved in the output:** Review became practical — instead of a dense paragraph that's hard to compare against the original resume, the table lets me decide row by row in a few minutes, hitting directly the 'the edit takes under 20 minutes' criterion from FL-01.

**What still didn't work:** The table made it easier to SPOT the repetition from V4, but didn't fix the repetition itself — the underlying text was still mechanical, it just became easier to reject row by row.

**What I'd try next:** Break the process into explicit steps (first map what's a real signal vs. a gap, then rewrite) to fix the repetition at the source, not just make it easier to catch.

## 7. Version 6 — Verification requirements (final prompt)

**Final prompt:** "You are a senior Data Quality Engineering recruiter... Context: ...My proof statement is: [paste]... Never invent or exaggerate experience. Do this in steps: (1) list the 6-8 core keywords from the job posting; (2) for each one, say whether the resume has real evidence (match) or not (gap) — don't invent matches; (3) rewrite only the items with a real match...; (4) for real gaps, suggest how to name them honestly. Deliver as a table..."

**Output (summary):** Step 2 (match vs. gap) explicitly revealed that 'formalized data testing' is a real gap — and the AI suggested naming it honestly as 'structured data validation via SQL', instead of inventing it. Step 3 avoided V4/V5's mechanical repetition, because the instruction asked for sentence variation. The final table was ready for direct use.

**What changed in the prompt:** Added an explicit verification requirement: for each keyword, the AI must first state whether there's real evidence in the resume (match) or not (gap) — without inventing matches — before rewriting anything. The step-by-step structure was just the mechanism used to apply that verification; the tested layer is 'verification requirements'.

**What actually improved in the output:** Separating 'mapping' from 'rewriting' eliminated the invention risk (which nearly showed up in V2) and solved V4/V5's mechanical repetition, because the rewrite now started from a verified mapping, not a single generic example.

**What still didn't work:** Running this side by side with ChatGPT (see comparison below) made it clear that Claude didn't flag when a field in the prompt itself (the proof statement) had been left with the empty placeholder — an execution error on my part that a real recruiter would notice right away. Claude also classified keywords less granularly than ChatGPT.

**What I'd try next:** Add a verification step at the start of the prompt ("confirm all bracketed fields are filled in before proceeding") to catch this kind of execution error, and ask the AI to classify the source of each piece of evidence (professional, project, academic) for more granularity.

## 8. Cross-model comparison — Claude vs. ChatGPT

The Version 6 final prompt was run on both models, with the same resume. Honest, point-by-point comparison:

**Structure.** Nearly identical — both delivered the keyword table (match/gap) followed by the Original | Suggestion | Justification table, exactly as requested in step 4 of the prompt.

**Accuracy / the most important finding.** ChatGPT caught an error left in the prompt itself: the proof statement field hadn't actually been filled in (it was still the literal placeholder "[paste]"), and it flagged this explicitly, refusing to invent a match from an empty field. Claude, on the same run, did not flag this — it proceeded assuming the proof statement context was implied. This is a real behavioral difference, not a style difference.

**Depth of analysis.** ChatGPT separated out 8 keywords (including ETL/Data Pipelines and Data Modeling as distinct items) and classified evidence by source — professional, project, academic, technical. Claude, in V6, treated the resume in a more unified way, with less granularity in that breakdown.

**Initiative beyond the requested scope.** ChatGPT proactively suggested rewriting the "Objective" line at the top of the resume (from "Jr. role in Data Analysis and Operations" to something like "Data Quality Engineer Jr. | Data Quality, SQL and Python") — a structural change the prompt didn't explicitly request, but one that is strategically relevant. Claude did not touch that line.

**Common ground (the most important point for the project's honesty criterion).** Both models converged exactly where it matters most: neither invented ETL experience or automated data testing experience. Both named these gaps with the same honest language ("data extraction and processing", not "ETL pipeline"). This confirms that the verification-requirements technique (Version 6) works consistently across models — it isn't a behavior specific to just one.

**Weak point identified.** Claude was less rigorous about flagging the prompt's fill-in failure (the empty "[paste]" placeholder) — a real recruiter would notice this immediately, so ChatGPT's check here has real practical value, not just style.

---

## 9. Reusable prompt template (final version)

Generic version, with no personal data from Fabiana, applicable by anyone to any resume + job pair:

You are a senior recruiter specialized in [JOB FIELD] evaluating resumes for the [ROLE] position at [COMPANY], whose stated focus is [JOB KEYWORDS].

Context: this resume supports the following statement about me: "[YOUR PROOF STATEMENT / VALUE SUMMARY]". It's essential that no adjustment invents or exaggerates experience — every sentence must stay 100% true.

Do this in steps: (1) list the 6-8 core keywords from the job posting; (2) for each one, say whether my resume has real evidence (match) or not (gap) — don't invent matches; (3) rewrite only the items with a real match, following the pattern action + result + explicit connection to the keyword, varying sentence structure; (4) for real gaps, suggest how to name them honestly, without pretending to have experience that doesn't exist.

Deliver as a table: Original | Suggestion | Justification (max 20 words).

[PASTE YOUR RESUME HERE]

## 10. Summary of techniques applied

V1 — None (baseline): generic advice, no use of the real resume.

V2 — Defined target audience: prioritized feedback, recruiter vocabulary.

V3 — Real context: connection to the proof statement; gaps named honestly.

V4 — Examples (few-shot): format consistency; risk of mechanical repetition.

V5 — Output format: reviewable table in minutes, not a dense paragraph.

V6 — Verification requirements: separated mapping from rewriting; eliminated the invention risk.