

The Prompt Ladder

Fabiana Alves Ferreira

This exercise uses a real mistake that happened in this very working session as the baseline — not a made-up example. It's exactly the kind of vague, context-free prompt this task asks you to recognize in your own history.

0. The weak task (baseline)

Real context: while reviewing a piece of already-submitted work (the FL-02 Prompt Iteration Log), I asked the AI to confirm whether all requirements had been met — without pasting the exact assignment text, relying instead on the conversation's memory.

Prompt: "Let's make sure we get an excellent grade on this task." / "Check if we met all the requirements for this task." (no assignment text attached)

Output (summary): The AI produced a confident-looking checklist — four items, all checked off — concluding "nothing is missing... you can post this with confidence." The problem: the four criteria used in that check belonged to a different assignment ("The Prompt Ladder"), not the one actually being evaluated ("Instruction Fundamentals on Real Tasks v2"). The confirmation looked rigorous, but it was checking the wrong thing.

This is the "slightly embarrassing" baseline result the task asks for — and it's real, not fabricated.

1. Version 1 — Clearer objective

Layer added: Clearer objective

Prompt: "Confirm, requirement by requirement, whether this document meets ALL the approval criteria stated in the assignment brief — not a general impression, a literal check of each criterion." (still no assignment text attached)

Output (summary): The response sounded even more rigorous — more assertive language, a more formal checklist — but it still evaluated against the wrong criteria, because the source of truth (the real assignment text) was still missing from the conversation.

What changed in the prompt: Asked for an explicit, requirement-by-requirement confirmation instead of a general impression.

What actually improved in the output: Nothing, in practice — the tone got more confident, but the basis for comparison was still wrong.

What still didn't work: This made things worse: a clearer objective, without the right source, just produces a wrong answer with MORE confidence — which is more dangerous than a vague answer, because it's easier to believe.

What I'd try next: Stop asking for a more rigorous tone and fix the actual problem: the AI didn't have the assignment text. Attach the real assignment.

2. Version 2 — Real context

Layer added: Real context

Prompt: Same confirmation request, but now with the full, literal text of the assignment pasted into the conversation, instead of relying on memory of a similar-looking task discussed earlier.

Output (summary): For the first time, the AI could actually compare the document against the correct assignment — and caught that two different Week 2 assignments had been mixed up with each other.

What changed in the prompt: Attached the complete, literal assignment text instead of relying on the conversation's memory.

What actually improved in the output: The check was now against the right source — the mix-up between two different assignments was only caught because the real text was there to compare against.

What still didn't work: Even with the right assignment text in hand, the response still made claims like "this matches criterion X" without pointing to the exact sentence in the document that proves it — it was still possible to claim a match without proving it.

What I'd try next: Require that every compliance claim come with an exact quote from the document that backs it up.

3. Version 3 — Verification requirements

Layer added: Verification requirements

Prompt: "For each criterion, quote the exact sentence from the assignment AND the exact passage from the document that satisfies it. If there's no passage that proves it, mark it as not met — don't assume."

Output (summary): The response stopped making loose claims. Every check or fail now came with two side-by-side quotes — the assignment's requirement and the document's evidence — making it impossible to confirm something without pointing to exactly where it happens.

What changed in the prompt: Added the requirement to quote, for each criterion, both the assignment's wording and the document passage that satisfies it — no quote, no credit.

What actually improved in the output: The check stopped being an opinion and became a traceable audit — anyone reading it can verify each claim themselves, instead of trusting the AI blindly.

What still didn't work: With the quotes embedded in running prose, it took effort to quickly scan which criteria passed and which didn't.

What I'd try next: Ask for a specific output format — a table — to make the audit fast to scan visually.

4. Version 4 — Specified output format

Layer added: Specified output format

Prompt: "Deliver the check as a table with four columns: Assignment criterion (quote) | Evidence in document (quote) | Status (Met / Not met / Partial) | Note."

Output (summary): A clean table, scannable in seconds — each criterion on its own row, the assignment quote right next to the document quote, with status visible at a glance.

What changed in the prompt: Specified the output format: a four-column table, with quotes side by side and an explicit status.

What actually improved in the output: The check became fast to review — you can see at a glance which criteria are actually met, without reading a paragraph of justification for each one.

What still didn't work: The table covers the criteria I know exist well — but it doesn't protect against the original risk: I could still be running the table against the wrong assignment again without noticing, and the table would look just as organized and convincing either way.

What I'd try next: Ask the AI to explicitly state any assumption or ambiguity before concluding — including which assignment/version is being used as the reference.

5. Version 5 — Stated assumptions (final prompt)

Layer added: Stated assumptions

Prompt: "Before filling in the table, explicitly state: (1) which assignment/version you're using as the reference, by name or code; (2) any ambiguity about which assignment applies to this document; (3) any assumption you're making due to missing information. Only then, fill in the verification table."

Output (summary): The response now opens by explicitly naming the assignment being used as the reference ("I'm evaluating this document against the FL-02 brief, code X, pasted into this conversation") before any checking happens — which would have prevented the original mistake, because the confusion between the two assignments would have been stated out loud before the conclusion was reached.

What changed in the prompt: Added the requirement to state, before concluding, which assignment is being used and any ambiguity or assumption involved.

What actually improved in the output: The exact type of error that caused this whole exercise — mixing up two similar assignments — is now visible BEFORE the conclusion, not after the damage is done. The reader can correct the AI before trusting the answer.

What still didn't work: This still depends on me actually reading the stated-assumptions section before accepting the answer — if I skip straight to the table, the safeguard doesn't work.

What I'd try next: Nothing further for this specific prompt — the final version already fixes the root cause of the original error. For general use, the next improvement would be requiring this same assumptions statement on any verification task, not just this one.

Reusable final prompt

Generic version, with no personal data, applicable by anyone to any compliance check against a written set of requirements:

You are auditing a document against a written set of requirements.

Step 1 — Before evaluating anything, state explicitly: (a) which requirements document/version you are using as the reference, by name; (b) any ambiguity about which requirements apply to this document; (c) any assumption you are making due to missing information.

Step 2 — For each requirement, quote the exact requirement text, then quote the exact passage in the document that satisfies it. If no passage satisfies it, mark it as not met — do not assume.

Step 3 — Deliver the result as a table: Requirement (quote) | Evidence in document (quote) | Status (Met / Not met / Partial) | Note.

Do not conclude the document is compliant unless every requirement has a quoted piece of evidence.

Summary of techniques applied

V1 — Clearer objective: more confident tone, without fixing the real problem — made things worse.

V2 — Real context: revealed the mix-up between two assignments, by comparing against the right brief.

V3 — Verification requirements: turned opinion into a traceable audit, with mandatory quotes.

V4 — Specified output format: scannable table, fast to review.

V5 — Stated assumptions: exposed which assignment is being used before the conclusion — the actual root-cause fix.